

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn import linear_model
```

*# Linear regression*

*#1. load Data*

```
from sklearn.datasets import load_boston
boston_dataset = load_boston()
boston_dataset =
pd.DataFrame(boston_dataset.data, columns=boston_dataset.feature_names)
```

C:\Users\Lenovo\AppData\Roaming\Python\Python310\site-packages\sklearn\utils\deprecation.py:87: FutureWarning: Function load\_boston is deprecated; `load\_boston` is deprecated in 1.0 and will be removed in 1.2.

The Boston housing prices dataset has an ethical problem. You can refer to the documentation of this function for further details.

The scikit-learn maintainers therefore strongly discourage the use of this dataset unless the purpose of the code is to study and educate about ethical issues in data science and machine learning.

In this special case, you can fetch the dataset from the original source::

```
import pandas as pd
import numpy as np

data_url = "http://lib.stat.cmu.edu/datasets/boston"
raw_df = pd.read_csv(data_url, sep="\s+", skiprows=22,
header=None)
data = np.hstack([raw_df.values[::2, :],
raw_df.values[1::2, :2]])
target = raw_df.values[1::2, 2]
```

Alternative datasets include the California housing dataset (i.e. :func:`~sklearn.datasets.fetch\_california\_housing`) and the Ames housing dataset. You can load the datasets as follows::

```
from sklearn.datasets import fetch_california_housing
```

```

housing = fetch_california_housing()

for the California housing dataset and::

    from sklearn.datasets import fetch_openml
    housing = fetch_openml(name="house_prices", as_frame=True)

for the Ames housing dataset.
warnings.warn(msg, category=FutureWarning)

```

```

#2. create Dataframe
#X,y= load_boston(return_X_y=True)
#X =pd.DataFrame(X)
#X.head()

```

C:\Users\Lenovo\AppData\Roaming\Python\Python310\site-packages\sklearn\utils\deprecation.py:87: FutureWarning: Function load\_boston is deprecated; `load\_boston` is deprecated in 1.0 and will be removed in 1.2.

The Boston housing prices dataset has an ethical problem. You can refer to the documentation of this function for further details.

The scikit-learn maintainers therefore strongly discourage the use of this dataset unless the purpose of the code is to study and educate about ethical issues in data science and machine learning.

In this special case, you can fetch the dataset from the original source::

```

import pandas as pd
import numpy as np

data_url = "http://lib.stat.cmu.edu/datasets/boston"
raw_df = pd.read_csv(data_url, sep="\s+", skiprows=22,
header=None)
data = np.hstack([raw_df.values[::2, :],
raw_df.values[1::2, :2]])
target = raw_df.values[1::2, 2]

```

Alternative datasets include the California housing dataset (i.e. :func:`~sklearn.datasets.fetch\_california\_housing`) and the Ames housing dataset. You can load the datasets as follows::

```

from sklearn.datasets import fetch_california_housing
housing = fetch_california_housing()

```

for the California housing dataset and::

```
from sklearn.datasets import fetch_openml
housing = fetch_openml(name="house_prices", as_frame=True)
```

```
for the Ames housing dataset.
warnings.warn(msg, category=FutureWarning)
```

	0	1	2	3	4	5	6	7	8	9
10 \										
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0
15.3										
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0
17.8										
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0
17.8										
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0
18.7										
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0
18.7										

	11	12
0	396.90	4.98
1	396.90	9.14
2	392.83	4.03
3	394.63	2.94
4	396.90	5.33

```
booston_dataset.head()
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX
\										
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0

	PTRATIO	B	LSTAT
0	15.3	396.90	4.98
1	17.8	396.90	9.14
2	17.8	392.83	4.03
3	18.7	394.63	2.94
4	18.7	396.90	5.33

```

booston_dataset.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 13 columns):
 #   Column      Non-Null Count  Dtype  
---  -
 0   CRIM        506 non-null    float64
 1   ZN          506 non-null    float64
 2   INDUS       506 non-null    float64
 3   CHAS        506 non-null    float64
 4   NOX         506 non-null    float64
 5   RM          506 non-null    float64
 6   AGE         506 non-null    float64
 7   DIS         506 non-null    float64
 8   RAD         506 non-null    float64
 9   TAX         506 non-null    float64
10  PTRATIO     506 non-null    float64
11  B           506 non-null    float64
12  LSTAT       506 non-null    float64
dtypes: float64(13)
memory usage: 51.5 KB

```

#### #4. create X and Y

```
X = booston_dataset
```

```
kashti = sns.load_dataset('titanic')
```

```
kashti
```

	survived	pclass	sex	age	sibsp	parch	fare	embarked
class \								
0	0	3	male	22.0	1	0	7.2500	S
Third								
1	1	1	female	38.0	1	0	71.2833	C
First								
2	1	3	female	26.0	0	0	7.9250	S
Third								
3	1	1	female	35.0	1	0	53.1000	S
First								
4	0	3	male	35.0	0	0	8.0500	S
Third								
..	...	...	...	...	...	...	...	...
...								
886	0	2	male	27.0	0	0	13.0000	S
Second								
887	1	1	female	19.0	0	0	30.0000	S
First								
888	0	3	female	NaN	1	2	23.4500	S
Third								
889	1	1	male	26.0	0	0	30.0000	C

```

First
890      0      3   male  32.0      0      0   7.7500      Q
Third

```

```

      who  adult_male  deck  embark_town  alive  alone
0      man      True  NaN  Southampton    no  False
1  woman      False   C    Cherbourg    yes  False
2  woman      False  NaN  Southampton    yes  True
3  woman      False   C    Southampton    yes  False
4      man      True  NaN  Southampton    no  True
..      ...      ...  ...      ...      ...
886   man      True  NaN  Southampton    no  True
887 woman      False   B    Southampton    yes  True
888 woman      False  NaN  Southampton    no  False
889   man      True   C    Cherbourg    yes  True
890   man      True  NaN  Queenstown    no  True

```

[891 rows x 15 columns]

```
kashti.head()
```

```

      survived  pclass      sex   age  sibsp  parch      fare embarked
class \
0      0      3   male  22.0      1      0   7.2500      S
Third
1      1      1  female  38.0      1      0  71.2833      C
First
2      1      3  female  26.0      0      0   7.9250      S
Third
3      1      1  female  35.0      1      0  53.1000      S
First
4      0      3   male  35.0      0      0   8.0500      S
Third

```

```

      who  adult_male  deck  embark_town  alive  alone
0      man      True  NaN  Southampton    no  False
1  woman      False   C    Cherbourg    yes  False
2  woman      False  NaN  Southampton    yes  True
3  woman      False   C    Southampton    yes  False
4      man      True  NaN  Southampton    no  True

```

```

# remove all null
#kashti.dropna(inplace=True)
#kashti

```

```

# remove null only from age col
kashti.dropna(subset= ['age'],inplace=True)
kashti

```

```

      survived  pclass      sex   age  sibsp  parch      fare embarked
class \

```

0	0	3	male	22.0	1	0	7.2500	S
Third								
1	1	1	female	38.0	1	0	71.2833	C
First								
2	1	3	female	26.0	0	0	7.9250	S
Third								
3	1	1	female	35.0	1	0	53.1000	S
First								
4	0	3	male	35.0	0	0	8.0500	S
Third								
..	...	...	...	...	...	...	...	...
...								
885	0	3	female	39.0	0	5	29.1250	Q
Third								
886	0	2	male	27.0	0	0	13.0000	S
Second								
887	1	1	female	19.0	0	0	30.0000	S
First								
889	1	1	male	26.0	0	0	30.0000	C
First								
890	0	3	male	32.0	0	0	7.7500	Q
Third								

	who	adult_male	deck	embark_town	alive	alone
0	man	True	NaN	Southampton	no	False
1	woman	False	C	Cherbourg	yes	False
2	woman	False	NaN	Southampton	yes	True
3	woman	False	C	Southampton	yes	False
4	man	True	NaN	Southampton	no	True
..	...	...	...	...	...	...
885	woman	False	NaN	Queenstown	no	False
886	man	True	NaN	Southampton	no	True
887	woman	False	B	Southampton	yes	True
889	man	True	C	Cherbourg	yes	True
890	man	True	NaN	Queenstown	no	True

[714 rows x 15 columns]

*# define X and y*

```
X= kashti[['age']]
y= kashti[['fare']]
```

*# check null value in data sey*

```
X.isnull().sum()
y.isnull().sum()
```

```
fare    0
dtype: int64
```

```
X.shape
```

```
(714, 1)
```

```
y.shape
```

```
(714, 1)
```

```
kashti.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 714 entries, 0 to 890
```

```
Data columns (total 15 columns):
```

#	Column	Non-Null Count	Dtype
0	survived	714 non-null	int64
1	pclass	714 non-null	int64
2	sex	714 non-null	object
3	age	714 non-null	float64
4	sibsp	714 non-null	int64
5	parch	714 non-null	int64
6	fare	714 non-null	float64
7	embarked	712 non-null	object
8	class	714 non-null	category
9	who	714 non-null	object
10	adult_male	714 non-null	bool
11	deck	184 non-null	category
12	embark_town	712 non-null	object
13	alive	714 non-null	object
14	alone	714 non-null	bool

```
dtypes: bool(2), category(2), float64(2), int64(4), object(5)
```

```
memory usage: 70.2+ KB
```

```
#linear Regersion
```

```
from sklearn.linear_model import LinearRegression
```

```
from sklearn.model_selection import train_test_split
```

```
X_train ,X_test,y_train ,y_test =
```

```
train_test_split(X,y ,test_size=0.2,random_state=0)
```

```
# train the model
```

```
model =LinearRegression()
```

```
model.fit(X_train,y_train)
```

```
LinearRegression()
```

```
# test the model
```

```
y_pred =model.predict(X_test)
```

```
# metrics to check the perfoermance of the model
```

```
from sklearn.metrics import mean_squared_error
```

```

y_pred =model.predict(X_test)
mean_squared_error(y_test,y_pred)

2128.1892760608985

from sklearn.metrics import
r2_score,mean_absolute_error,mean_absolute_error
import numpy as np
r2_score(y_test ,y_pred)

0.015800553474554446

mae = mean_absolute_error(y_test ,y_pred)
mse = mean_squared_error(y_test ,y_pred)
rmse =np.sqrt(mse)
r2=r2_score(y_test ,y_pred)

print('Mean Absolute Error ',mae)
print('Mean Squared Errorr',mse)
print('Root Mean Squared Error',rmse)
print('R2 Score',r2 )

Mean Absolute Error  31.184262095057967
Mean Squared Errorr 2128.1892760608985
Root Mean Squared Error 46.13230187255887
R2 Score 0.015800553474554446

mse = mean_squared_error(y_test ,y_pred)
mse

2128.1892760608985

rmse =np.sqrt(mse)
rmse


46.13230187255887

#plot
kashti.plot(kind='scatter', x='age',y='fare',color='red')

<AxesSubplot:xlabel='age', ylabel='fare'>

```





Mean Absolute Error 31.184262095057967

Mean Absolute Error (MAE): When we subtract the predicted values from the actual values, obtaining the errors, sum the absolute values of those errors and get their mean. This metric gives a notion of the overall error for each prediction of the model, the smaller (closer to 0) the better.

Mean Squared Error 2128.1892760608985

Mean Squared Error (MSE): It is similar to the MAE metric, but it squares the absolute values of the errors. Also, as with MAE, the smaller, or closer to 0, the better. The MSE value is squared so as to make large errors even larger. One thing to pay close attention to, it that it is usually a hard metric to interpret due to the size of its values and of the fact that they aren't in the same scale of the data.

Root Mean Squared Error 46.13230187255887

Root Mean Squared Error (RMSE): Tries to solve the interpretation problem raised with the MSE by getting the square root of its final value, so as to scale it back to the same units of the data. It is easier to interpret and good when we need to display or show the actual value of the data with the error. It shows how much the data may vary, so, if we have an RMSE of 4.35, our model can make an error either because it added 4.35 to the actual value, or needed 4.35 to get to the actual value. The closer to 0, the better as well.

R2 Score 0.015800553474554446

R-Squared ( $R^2$  or the coefficient of determination) is a statistical measure in a regression model that determines the proportion of variance in the dependent variable that can be explained by the independent variable. In other words, r-squared shows how well the data fit the regression model (the goodness of fit).